

## 7.1. string — Common string operations

---

**Source code:** [Lib/string.py](#)

---

The `string` module contains a number of useful constants and classes, as well as some deprecated legacy functions that are also available as methods on strings. In addition, Python's built-in string classes support the sequence type methods described in the [Sequence Types — str, unicode, list, tuple, bytearray, buffer, xrange](#) section, and also the string-specific methods described in the [String Methods](#) section. To output formatted strings use template strings or the `%` operator described in the [String Formatting Operations](#) section. Also, see the `re` module for string functions based on regular expressions.

### 7.1.1. String constants

---

The constants defined in this module are:

`string.ascii_letters`

The concatenation of the `ascii_lowercase` and `ascii_uppercase` constants described below. This value is not locale-dependent.

`string.ascii_lowercase`

The lowercase letters `'abcdefghijklmnopqrstuvwxyz'`. This value is not locale-dependent and will not change.

`string.ascii_uppercase`

The uppercase letters `'ABCDEFGHIJKLMNOPQRSTUVWXYZ'`. This value is not locale-dependent and will not change.

`string.digits`

The string `'0123456789'`.

`string.hexdigits`

The string `'0123456789abcdefABCDEF'`.

`string.letters`

The concatenation of the strings `lowercase` and `uppercase` described below. The specific value is locale-dependent, and will be updated when `locale.setlocale()` is called.

`string.lowercase`

A string containing all the characters that are considered lowercase letters. On most systems this is the string `'abcdefghijklmnopqrstuvwxyz'`. The specific value is locale-dependent, and will be updated when `locale.setlocale()` is called.

`string.octdigits`

The string `'01234567'`.

`string.punctuation`

String of ASCII characters which are considered punctuation characters in the `c` locale.

`string.printable`

String of characters which are considered printable. This is a combination of `digits`, `letters`, `punctuation`, and `whitespace`.

`string.uppercase`

A string containing all the characters that are considered uppercase letters. On most systems this is the string `'ABCDEFGHIJKLMNOPQRSTUVWXYZ'`. The specific value is locale-dependent, and will be updated when `locale.setlocale()` is called.

`string.whitespace`

A string containing all characters that are considered whitespace. On most systems this includes the characters space, tab, linefeed, return, formfeed, and vertical tab.

## 7.1.2. Custom String Formatting

---

*New in version 2.6.*

The built-in `str` and `unicode` classes provide the ability to do complex variable substitutions and value formatting via the `str.format()` method described in [PEP 3101](#). The `Formatter` class in the `string` module allows you to create and customize your own string formatting behaviors using the same implementation as the built-in `format()` method.

`class string.Formatter`

The `Formatter` class has the following public methods:

**`format(format_string, *args, **kwargs)`**

The primary API method. It takes a format string and an arbitrary set of positional and keyword arguments. It is just a wrapper that calls `vformat()`.

**`vformat(format_string, args, kwargs)`**

This function does the actual work of formatting. It is exposed as a separate function for cases where you want to pass in a predefined dictionary of arguments, rather than unpacking and repacking the dictionary as individual arguments using the `*args` and `**kwargs` syntax. `vformat()` does the work of breaking up the format string into character data and replacement fields. It calls the various methods described below.

In addition, the `Formatter` defines a number of methods that are intended to be replaced by subclasses:

**`parse(format_string)`**

Loop over the `format_string` and return an iterable of tuples (`literal_text`, `field_name`, `format_spec`, `conversion`). This is used by `vformat()` to break the string into either literal text, or replacement fields.

The values in the tuple conceptually represent a span of literal text followed by a single replacement field. If there is no literal text (which can happen if two replacement fields

occur consecutively), then *literal\_text* will be a zero-length string. If there is no replacement field, then the values of *field\_name*, *format\_spec* and *conversion* will be `None`.

#### **get\_field(*field\_name*, *args*, *kwargs*)**

Given *field\_name* as returned by `parse()` (see above), convert it to an object to be formatted. Returns a tuple (obj, used\_key). The default version takes strings of the form defined in **PEP 3101**, such as “0[name]” or “label.title”. *args* and *kwargs* are as passed in to `vformat()`. The return value *used\_key* has the same meaning as the *key* parameter to `get_value()`.

#### **get\_value(*key*, *args*, *kwargs*)**

Retrieve a given field value. The *key* argument will be either an integer or a string. If it is an integer, it represents the index of the positional argument in *args*; if it is a string, then it represents a named argument in *kwargs*.

The *args* parameter is set to the list of positional arguments to `vformat()`, and the *kwargs* parameter is set to the dictionary of keyword arguments.

For compound field names, these functions are only called for the first component of the field name; Subsequent components are handled through normal attribute and indexing operations.

So for example, the field expression ‘0.name’ would cause `get_value()` to be called with a *key* argument of 0. The `name` attribute will be looked up after `get_value()` returns by calling the built-in `getattr()` function.

If the index or keyword refers to an item that does not exist, then an `IndexError` or `KeyError` should be raised.

#### **check\_unused\_args(*used\_args*, *args*, *kwargs*)**

Implement checking for unused arguments if desired. The arguments to this function is the set of all argument keys that were actually referred to in the format string (integers for positional arguments, and strings for named arguments), and a reference to the *args* and *kwargs* that was passed to `vformat`. The set of unused args can be calculated from these parameters. `check_unused_args()` is assumed to raise an exception if the check fails.

#### **format\_field(*value*, *format\_spec*)**

`format_field()` simply calls the global `format()` built-in. The method is provided so that subclasses can override it.

#### **convert\_field(*value*, *conversion*)**

Converts the value (returned by `get_field()`) given a conversion type (as in the tuple returned by the `parse()` method). The default version understands ‘s’ (str), ‘r’ (repr) and ‘a’ (ascii) conversion types.

## 7.1.3. Format String Syntax

---

The `str.format()` method and the `Formatter` class share the same syntax for format strings (although in the case of `Formatter`, subclasses can define their own format string syntax).

Format strings contain “replacement fields” surrounded by curly braces `{}`. Anything that is not contained in braces is considered literal text, which is copied unchanged to the output. If you need to include a brace character in the literal text, it can be escaped by doubling: `{{` and `}}`.

The grammar for a replacement field is as follows:

---

```
replacement_field ::= "{" [field_name] ["!" conversion] [":" format_spec] "}"
field_name        ::= arg_name ("." attribute_name | "[" element_index "]")*
arg_name          ::= [identifier | integer]
attribute_name    ::= identifier
element_index     ::= integer | index_string
index_string      ::= <any source character except "]"> +
conversion        ::= "r" | "s"
format_spec       ::= <described in the next section>
```

---

In less formal terms, the replacement field can start with a *field\_name* that specifies the object whose value is to be formatted and inserted into the output instead of the replacement field. The *field\_name* is optionally followed by a *conversion* field, which is preceded by an exclamation point `!`, and a *format\_spec*, which is preceded by a colon `:`. These specify a non-default format for the replacement value.

See also the [Format Specification Mini-Language](#) section.

The *field\_name* itself begins with an *arg\_name* that is either a number or a keyword. If it's a number, it refers to a positional argument, and if it's a keyword, it refers to a named keyword argument. If the numerical *arg\_names* in a format string are 0, 1, 2, ... in sequence, they can all be omitted (not just some) and the numbers 0, 1, 2, ... will be automatically inserted in that order. Because *arg\_name* is not quote-delimited, it is not possible to specify arbitrary dictionary keys (e.g., the strings `'10'` or `':-:]'`) within a format string. The *arg\_name* can be followed by any number of index or attribute expressions. An expression of the form `.'name'` selects the named attribute using `getattr()`, while an expression of the form `'[index]'` does an index lookup using `__getitem__()`.

*Changed in version 2.7:* The positional argument specifiers can be omitted for `str.format()` and `unicode.format()`, so `'{} {}'` is equivalent to `'{0} {1}'`, `u'{} {}'` is equivalent to `u'{0} {1}'`.

Some simple format string examples:

---

```
"First, thou shalt count to {0}" # References first positional argument
"Bring me a {}"                 # Implicitly references the first positional argument
"From {} to {}".format(0, 1)    # Same as "From {0} to {1}"
"My quest is {name}"            # References keyword argument 'name'
"Weight in tons {0.weight}"     # 'weight' attribute of first positional arg
"Units destroyed: {players[0]}" # First element of keyword argument 'players'.
```

---

The *conversion* field causes a type coercion before formatting. Normally, the job of formatting a value is done by the `__format__()` method of the value itself. However, in some cases it is desirable to force a type to be formatted as a string, overriding its own definition of formatting. By converting the value to a string before calling `__format__()`, the normal formatting logic is bypassed.

Two conversion flags are currently supported: `'!s'` which calls `str()` on the value, and `'!r'` which calls `repr()`.

Some examples:

---

```
"Harold's a clever {0!s}"      # Calls str() on the argument first
"Bring out the holy {name!r}"  # Calls repr() on the argument first
```

---

The `format_spec` field contains a specification of how the value should be presented, including such details as field width, alignment, padding, decimal precision and so on. Each value type can define its own “formatting mini-language” or interpretation of the `format_spec`.

Most built-in types support a common formatting mini-language, which is described in the next section.

A `format_spec` field can also include nested replacement fields within it. These nested replacement fields may contain a field name, conversion flag and format specification, but deeper nesting is not allowed. The replacement fields within the `format_spec` are substituted before the `format_spec` string is interpreted. This allows the formatting of a value to be dynamically specified.

See the [Format examples](#) section for some examples.

### 7.1.3.1. Format Specification Mini-Language

“Format specifications” are used within replacement fields contained within a format string to define how individual values are presented (see [Format String Syntax](#)). They can also be passed directly to the built-in `format()` function. Each formattable type may define how the format specification is to be interpreted.

Most built-in types implement the following options for format specifications, although some of the formatting options are only supported by the numeric types.

A general convention is that an empty format string ("" ) produces the same result as if you had called `str()` on the value. A non-empty format string typically modifies the result.

The general form of a *standard format specifier* is:

---

```
format_spec ::= [[fill]align][sign][#][0][width][,][.precision][type]
fill         ::= <any character>
align        ::= "<" | ">" | "=" | "^"
sign         ::= "+" | "-" | " "
width        ::= integer
precision    ::= integer
type         ::= "b" | "c" | "d" | "e" | "E" | "f" | "F" | "g" | "G" | "n" | "o" | "s" | "x" | "X"
```

---

If a valid `align` value is specified, it can be preceded by a `fill` character that can be any character and defaults to a space if omitted. It is not possible to use a literal curly brace (“{” or “}”) as the `fill` character when using the `str.format()` method. However, it is possible to insert a curly brace with a nested replacement field. This limitation doesn’t affect the `format()` function.

The meaning of the various alignment options is as follows:

| Option | Meaning  |
|--------|--|
| '<'    | Forces the field to be left-aligned within the available space (this is the default for most objects).   |
| '>'    | Forces the field to be right-aligned within the available space (this is the default for numbers).   |
| '='    | Forces the padding to be placed after the sign (if any) but before the digits. This is used for printing fields in the form '+000000120'. This alignment option is only valid for numeric types. It becomes the default when '0' immediately precedes the field width. |
| '^'    | Forces the field to be centered within the available space.  |

Note that unless a minimum field width is defined, the field width will always be the same size as the data to fill it, so that the alignment option has no meaning in this case.

The *sign* option is only valid for number types, and can be one of the following:

| Option | Meaning  |
|--------|--|
| '+'    | indicates that a sign should be used for both positive as well as negative numbers.                      |
| '-'    | indicates that a sign should be used only for negative numbers (this is the default behavior).           |
| space  | indicates that a leading space should be used on positive numbers, and a minus sign on negative numbers. |

The '#' option is only valid for integers, and only for binary, octal, or hexadecimal output. If present, it specifies that the output will be prefixed by '0b', '0o', or '0x', respectively.

The ',' option signals the use of a comma for a thousands separator. For a locale aware separator, use the 'n' integer presentation type instead.

*Changed in version 2.7:* Added the ',' option (see also [PEP 378](#)).

*width* is a decimal integer defining the minimum field width. If not specified, then the field width will be determined by the content.

When no explicit alignment is given, preceding the *width* field by a zero ('0') character enables sign-aware zero-padding for numeric types. This is equivalent to a *fill* character of '0' with an *alignment* type of '='.

The *precision* is a decimal number indicating how many digits should be displayed after the decimal point for a floating point value formatted with 'f' and 'F', or before and after the decimal point for a floating point value formatted with 'g' or 'G'. For non-number types the field indicates the maximum field size - in other words, how many characters will be used from the field content. The *precision* is not allowed for integer values.

Finally, the *type* determines how the data should be presented.

The available string presentation types are:

| Type | Meaning   |
|------|---|
| 's'  | String format. This is the default type for strings and may be omitted. |
| None | The same as 's'.  |

The available integer presentation types are:

| Type | Meaning  |
|------|--|
| 'b'  | Binary format. Outputs the number in base 2.   |
| 'c'  | Character. Converts the integer to the corresponding unicode character before printing.  |
| 'd'  | Decimal Integer. Outputs the number in base 10.  |
| 'o'  | Octal format. Outputs the number in base 8.  |
| 'x'  | Hex format. Outputs the number in base 16, using lower- case letters for the digits above 9.   |
| 'X'  | Hex format. Outputs the number in base 16, using upper- case letters for the digits above 9.   |
| 'n'  | Number. This is the same as 'd', except that it uses the current locale setting to insert the appropriate number separator characters. |
| None | The same as 'd'.   |

In addition to the above presentation types, integers can be formatted with the floating point presentation types listed below (except 'n' and None). When doing so, `float()` is used to convert the integer to a floating point number before formatting.

The available presentation types for floating point and decimal values are:

| Type | Meaning  |
|------|--|
| 'e'  | Exponent notation. Prints the number in scientific notation using the letter 'e' to indicate the exponent. The default precision is 6. |
| 'E'  | Exponent notation. Same as 'e' except it uses an upper case 'E' as the separator character.  |
| 'f'  | Fixed-point notation. Displays the number as a fixed-point number. The default precision is 6.   |
| 'F'  | Fixed point notation. Same as 'f'.   |



| Type | Meaning  |
|------|--|
| 'g'  | <p>General format. For a given precision <math>p \geq 1</math>, this rounds the number to <math>p</math> significant digits and then formats the result in either fixed-point format or in scientific notation, depending on its magnitude.</p> <p>The precise rules are as follows: suppose that the result formatted with presentation type 'e' and precision <math>p-1</math> would have exponent <math>exp</math>. Then if <math>-4 \leq exp &lt; p</math>, the number is formatted with presentation type 'f' and precision <math>p-1-exp</math>. Otherwise, the number is formatted with presentation type 'e' and precision <math>p-1</math>. In both cases insignificant trailing zeros are removed from the significand, and the decimal point is also removed if there are no remaining digits following it.</p> <p>Positive and negative infinity, positive and negative zero, and nans, are formatted as <code>inf</code>, <code>-inf</code>, <code>0</code>, <code>-0</code> and <code>nan</code> respectively, regardless of the precision.</p> <p>A precision of <code>0</code> is treated as equivalent to a precision of <code>1</code>. The default precision is <code>6</code>.</p> |
| 'G'  | General format. Same as 'g' except switches to 'E' if the number gets too large. The representations of infinity and NaN are uppercased, too.  |
| 'n'  | Number. This is the same as 'g', except that it uses the current locale setting to insert the appropriate number separator characters.   |
| '%'  | Percentage. Multiplies the number by 100 and displays in fixed ('f') format, followed by a percent sign.   |
| None | The same as 'g'.   |

### 7.1.3.2. Format examples

This section contains examples of the `str.format()` syntax and comparison with the old %-formatting.

In most of the cases the syntax is similar to the old %-formatting, with the addition of the `{}` and with `:` used instead of `%`. For example, `'%03.2f'` can be translated to `'{:03.2f}'`.

The new format syntax also supports new and different options, shown in the following examples.

Accessing arguments by position:

```
>>> '{0}, {1}, {2}'.format('a', 'b', 'c')
'a, b, c'
>>> '{}, {}, {}'.format('a', 'b', 'c') # 2.7+ only
'a, b, c'
>>> '{2}, {1}, {0}'.format('a', 'b', 'c')
'c, b, a'
>>> '{2}, {1}, {0}'.format(*'abc')      # unpacking argument sequence
'c, b, a'
>>> '{0}{1}{0}'.format('abra', 'cad')  # arguments' indices can be repeated
'abracadabra'
```



## Accessing arguments by name:

---

```
>>> 'Coordinates: {latitude}, {longitude}'.format(latitude='37.24N', longitude='-115.81W')
'Coordinates: 37.24N, -115.81W'
>>> coord = {'latitude': '37.24N', 'longitude': '-115.81W'}
>>> 'Coordinates: {latitude}, {longitude}'.format(**coord)
'Coordinates: 37.24N, -115.81W'
```

---

## Accessing arguments' attributes:

---

```
>>> c = 3-5j
>>> ('The complex number {0} is formed from the real part {0.real} '
... 'and the imaginary part {0.imag}').format(c)
'The complex number (3-5j) is formed from the real part 3.0 and the imaginary part -5.0.'
>>> class Point(object):
...     def __init__(self, x, y):
...         self.x, self.y = x, y
...     def __str__(self):
...         return 'Point({self.x}, {self.y})'.format(self=self)
...
>>> str(Point(4, 2))
'Point(4, 2)'
```

---

## Accessing arguments' items:

---

```
>>> coord = (3, 5)
>>> 'X: {0[0]}; Y: {0[1]}'.format(coord)
'X: 3; Y: 5'
```

---

## Replacing %s and %r:

---

```
>>> "repr() shows quotes: {!r}; str() doesn't: {!s}".format('test1', 'test2')
'repr() shows quotes: \'test1\'; str() doesn\'t: test2'
```

---

## Aligning the text and specifying a width:

---

```
>>> '{:<30}'.format('left aligned')
'left aligned'
>>> '{:>30}'.format('right aligned')
'right aligned'
>>> '{:^30}'.format('centered')
'centered'
>>> '{:*^30}'.format('centered') # use '*' as a fill char
'*****centered*****'
```

---

## Replacing %f, %-f, and % f and specifying a sign:

---

```
>>> '{:+f}; {:+f}'.format(3.14, -3.14) # show it always
'+3.140000; -3.140000'
>>> '{: f}; {: f}'.format(3.14, -3.14) # show a space for positive numbers
' 3.140000; -3.140000'
>>> '{:-f}; {:-f}'.format(3.14, -3.14) # show only the minus -- same as '{:f}; {:f}'
'3.140000; -3.140000'
```

---

## Replacing %x and %o and converting the value to different bases:

---

```
>>> # format also supports binary numbers
>>> "int: {0:d}; hex: {0:x}; oct: {0:o}; bin: {0:b}".format(42)
'int: 42; hex: 2a; oct: 52; bin: 101010'
>>> # with 0x, 0o, or 0b as prefix:
>>> "int: {0:d}; hex: {0:#x}; oct: {0:#o}; bin: {0:#b}".format(42)
'int: 42; hex: 0x2a; oct: 0o52; bin: 0b101010'
```

---

Using the comma as a thousands separator:

---

```
>>> '{:,}'.format(1234567890)
'1,234,567,890'
```

---

Expressing a percentage:

---

```
>>> points = 19.5
>>> total = 22
>>> 'Correct answers: {:.2%}'.format(points/total)
'Correct answers: 88.64%'
```

---

Using type-specific formatting:

---

```
>>> import datetime
>>> d = datetime.datetime(2010, 7, 4, 12, 15, 58)
>>> '{:%Y-%m-%d %H:%M:%S}'.format(d)
'2010-07-04 12:15:58'
```

---

Nesting arguments and more complex examples:

---

```
>>> for align, text in zip('<^>', ['left', 'center', 'right']):
...     '{0:{fill}{align}16}'.format(text, fill=align, align=align)
...
'left<<<<<<<<<<<<<'
'^^^^^center^^^^^'
'>>>>>>>>>>>>>right'
>>>
>>> octets = [192, 168, 0, 1]
>>> '{:02X}{:02X}{:02X}{:02X}'.format(*octets)
'C0A80001'
>>> int(_, 16)
3232235521
>>>
>>> width = 5
>>> for num in range(5,12):
...     for base in 'dXob':
...         print '{0:{width}{base}}'.format(num, base=base, width=width),
...     print
...
5      5      5      101
6      6      6      110
7      7      7      111
8      8      10     1000
9      9      11     1001
10     A      12     1010
11     B      13     1011
```

---

## 7.1.4. Template strings

*New in version 2.4.*

Templates provide simpler string substitutions as described in [PEP 292](#). Instead of the normal %-based substitutions, Templates support \$-based substitutions, using the following rules:

- `$$` is an escape; it is replaced with a single `$`.
- `$identifier` names a substitution placeholder matching a mapping key of `"identifier"`. By default, `"identifier"` must spell a Python identifier. The first non-identifier character after the `$` character terminates this placeholder specification.
- `${identifier}` is equivalent to `$identifier`. It is required when valid identifier characters follow the placeholder but are not part of the placeholder, such as `"${noun}ification"`.

Any other appearance of `$` in the string will result in a `ValueError` being raised.

The `string` module provides a `Template` class that implements these rules. The methods of `Template` are:

**`class string.Template(template)`**

The constructor takes a single argument which is the template string.

**`substitute(mapping[, **kws])`**

Performs the template substitution, returning a new string. *mapping* is any dictionary-like object with keys that match the placeholders in the template. Alternatively, you can provide keyword arguments, where the keywords are the placeholders. When both *mapping* and *kws* are given and there are duplicates, the placeholders from *kws* take precedence.

**`safe_substitute(mapping[, **kws])`**

Like `substitute()`, except that if placeholders are missing from *mapping* and *kws*, instead of raising a `KeyError` exception, the original placeholder will appear in the resulting string intact. Also, unlike with `substitute()`, any other appearances of the `$` will simply return `$` instead of raising `ValueError`.

While other exceptions may still occur, this method is called “safe” because it always tries to return a usable string instead of raising an exception. In another sense, `safe_substitute()` may be anything other than safe, since it will silently ignore malformed templates containing dangling delimiters, unmatched braces, or placeholders that are not valid Python identifiers.

`Template` instances also provide one public data attribute:

**`template`**

This is the object passed to the constructor’s *template* argument. In general, you shouldn’t change it, but read-only access is not enforced.

Here is an example of how to use a Template:

---

```
>>> from string import Template
>>> s = Template('$who likes $what')
>>> s.substitute(who='tim', what='kung pao')
```

```
'tim likes kung pao'
>>> d = dict(who='tim')
>>> Template('Give $who $100').substitute(d)
Traceback (most recent call last):
...
ValueError: Invalid placeholder in string: line 1, col 11
>>> Template('$who likes $what').substitute(d)
Traceback (most recent call last):
...
KeyError: 'what'
>>> Template('$who likes $what').safe_substitute(d)
'tim likes $what'
```

Advanced usage: you can derive subclasses of `Template` to customize the placeholder syntax, delimiter character, or the entire regular expression used to parse template strings. To do this, you can override these class attributes:

- *delimiter* – This is the literal string describing a placeholder introducing delimiter. The default value is `$`. Note that this should *not* be a regular expression, as the implementation will call `re.escape()` on this string as needed.
- *idpattern* – This is the regular expression describing the pattern for non-braced placeholders (the braces will be added automatically as appropriate). The default value is the regular expression `[_a-z][_a-z0-9]*`.

Alternatively, you can provide the entire regular expression pattern by overriding the class attribute *pattern*. If you do this, the value must be a regular expression object with four named capturing groups. The capturing groups correspond to the rules given above, along with the invalid placeholder rule:

- *escaped* – This group matches the escape sequence, e.g. `$$`, in the default pattern.
- *named* – This group matches the unbraced placeholder name; it should not include the delimiter in capturing group.
- *braced* – This group matches the brace enclosed placeholder name; it should not include either the delimiter or braces in the capturing group.
- *invalid* – This group matches any other delimiter pattern (usually a single delimiter), and it should appear last in the regular expression.

## 7.1.5. String functions

The following functions are available to operate on string and Unicode objects. They are not available as string methods.

`string.capwords(s[, sep])`

Split the argument into words using `str.split()`, capitalize each word using `str.capitalize()`, and join the capitalized words using `str.join()`. If the optional second argument *sep* is absent or `None`, runs of whitespace characters are replaced by a single space and leading and trailing whitespace are removed, otherwise *sep* is used to split and join the words.

`string.maketrans(from, to)`

Return a translation table suitable for passing to `translate()`, that will map each character in *from* into the character at the same position in *to*; *from* and *to* must have the same length.

**Note:** Don't use strings derived from `lowercase` and `uppercase` as arguments; in some locales, these don't have the same length. For case conversions, always use `str.lower()` and `str.upper()`.

## 7.1.6. Deprecated string functions

The following list of functions are also defined as methods of string and Unicode objects; see section [String Methods](#) for more information on those. You should consider these functions as deprecated, although they will not be removed until Python 3. The functions defined in this module are:

string.**atof**(s)

*Deprecated since version 2.0:* Use the `float()` built-in function.

Convert a string to a floating point number. The string must have the standard syntax for a floating point literal in Python, optionally preceded by a sign (+ or -). Note that this behaves identical to the built-in function `float()` when passed a string.

**Note:** When passing in a string, values for NaN and Infinity may be returned, depending on the underlying C library. The specific set of strings accepted which cause these values to be returned depends entirely on the C library and is known to vary.

string.**atoi**(s[, base])

*Deprecated since version 2.0:* Use the `int()` built-in function.

Convert string `s` to an integer in the given `base`. The string must consist of one or more digits, optionally preceded by a sign (+ or -). The `base` defaults to 10. If it is 0, a default base is chosen depending on the leading characters of the string (after stripping the sign): `0x` or `0X` means 16, `0` means 8, anything else means 10. If `base` is 16, a leading `0x` or `0X` is always accepted, though not required. This behaves identically to the built-in function `int()` when passed a string. (Also note: for a more flexible interpretation of numeric literals, use the built-in function `eval()`.)

string.**atol**(s[, base])

*Deprecated since version 2.0:* Use the `long()` built-in function.

Convert string `s` to a long integer in the given `base`. The string must consist of one or more digits, optionally preceded by a sign (+ or -). The `base` argument has the same meaning as for `atoi()`. A trailing `l` or `L` is not allowed, except if the base is 0. Note that when invoked without `base` or with `base` set to 10, this behaves identical to the built-in function `long()` when passed a string.

string.**capitalize**(word)

Return a copy of `word` with only its first character capitalized.

string.**expandtabs**(s[, tabsize])

Expand tabs in a string replacing them by one or more spaces, depending on the current column and the given tab size. The column number is reset to zero after each newline occurring in the string. This doesn't understand other non-printing characters or escape sequences. The tab size defaults to 8.

`string.find(s, sub[, start[, end]])`

Return the lowest index in *s* where the substring *sub* is found such that *sub* is wholly contained in *s*[*start*:*end*]. Return -1 on failure. Defaults for *start* and *end* and interpretation of negative values is the same as for slices.

`string.rfind(s, sub[, start[, end]])`

Like `find()` but find the highest index.

`string.index(s, sub[, start[, end]])`

Like `find()` but raise `ValueError` when the substring is not found.

`string.rindex(s, sub[, start[, end]])`

Like `rfind()` but raise `ValueError` when the substring is not found.

`string.count(s, sub[, start[, end]])`

Return the number of (non-overlapping) occurrences of substring *sub* in string *s*[*start*:*end*]. Defaults for *start* and *end* and interpretation of negative values are the same as for slices.

`string.lower(s)`

Return a copy of *s*, but with upper case letters converted to lower case.

`string.split(s[, sep[, maxsplit]])`

Return a list of the words of the string *s*. If the optional second argument *sep* is absent or `None`, the words are separated by arbitrary strings of whitespace characters (space, tab, newline, return, formfeed). If the second argument *sep* is present and not `None`, it specifies a string to be used as the word separator. The returned list will then have one more item than the number of non-overlapping occurrences of the separator in the string. If *maxsplit* is given, at most *maxsplit* number of splits occur, and the remainder of the string is returned as the final element of the list (thus, the list will have at most *maxsplit*+1 elements). If *maxsplit* is not specified or -1, then there is no limit on the number of splits (all possible splits are made).

The behavior of `split` on an empty string depends on the value of *sep*. If *sep* is not specified, or specified as `None`, the result will be an empty list. If *sep* is specified as any string, the result will be a list containing one element which is an empty string.

`string.rsplit(s[, sep[, maxsplit]])`

Return a list of the words of the string *s*, scanning *s* from the end. To all intents and purposes, the resulting list of words is the same as returned by `split()`, except when the optional third argument *maxsplit* is explicitly specified and nonzero. If *maxsplit* is given, at most *maxsplit* number of splits – the *rightmost* ones – occur, and the remainder of the string is returned as the first element of the list (thus, the list will have at most *maxsplit*+1 elements).

*New in version 2.4.*

`string.splitfields(s[, sep[, maxsplit]])`

This function behaves identically to `split()`. (In the past, `split()` was only used with one argument, while `splitfields()` was only used with two arguments.)

`string.join(words[, sep])`

Concatenate a list or tuple of words with intervening occurrences of `sep`. The default value for `sep` is a single space character. It is always true that `string.join(string.split(s, sep), sep)` equals `s`.

`string.joinfields(words[, sep])`

This function behaves identically to `join()`. (In the past, `join()` was only used with one argument, while `joinfields()` was only used with two arguments.) Note that there is no `joinfields()` method on string objects; use the `join()` method instead.

`string.lstrip(s[, chars])`

Return a copy of the string with leading characters removed. If `chars` is omitted or `None`, whitespace characters are removed. If given and not `None`, `chars` must be a string; the characters in the string will be stripped from the beginning of the string this method is called on.

*Changed in version 2.2.3:* The `chars` parameter was added. The `chars` parameter cannot be passed in earlier 2.2 versions.

`string.rstrip(s[, chars])`

Return a copy of the string with trailing characters removed. If `chars` is omitted or `None`, whitespace characters are removed. If given and not `None`, `chars` must be a string; the characters in the string will be stripped from the end of the string this method is called on.

*Changed in version 2.2.3:* The `chars` parameter was added. The `chars` parameter cannot be passed in earlier 2.2 versions.

`string.strip(s[, chars])`

Return a copy of the string with leading and trailing characters removed. If `chars` is omitted or `None`, whitespace characters are removed. If given and not `None`, `chars` must be a string; the characters in the string will be stripped from the both ends of the string this method is called on.

*Changed in version 2.2.3:* The `chars` parameter was added. The `chars` parameter cannot be passed in earlier 2.2 versions.

`string.swapcase(s)`

Return a copy of `s`, but with lower case letters converted to upper case and vice versa.

`string.translate(s, table[, deletechars])`

Delete all characters from `s` that are in `deletechars` (if present), and then translate the characters using `table`, which must be a 256-character string giving the translation for each



character value, indexed by its ordinal. If *table* is `None`, then only the character deletion step is performed.

`string.upper(s)`

Return a copy of *s*, but with lower case letters converted to upper case.

`string.ljust(s, width[, fillchar])`

`string.rjust(s, width[, fillchar])`

`string.center(s, width[, fillchar])`

These functions respectively left-justify, right-justify and center a string in a field of given width. They return a string that is at least *width* characters wide, created by padding the string *s* with the character *fillchar* (default is a space) until the given width on the right, left or both sides. The string is never truncated.

`string.zfill(s, width)`

Pad a numeric string *s* on the left with zero digits until the given *width* is reached. Strings starting with a sign are handled correctly.

`string.replace(s, old, new[, maxreplace])`

Return a copy of string *s* with all occurrences of substring *old* replaced by *new*. If the optional argument *maxreplace* is given, the first *maxreplace* occurrences are replaced.